

Research Interests

- Scalable compiler directed workload analysis
- Hardware software co-design for specialized architectures
- Core micro-architecture with a focus on the cache memory hierarchy

Academic

05/13 – 11/16 **PhD in Computing Science, Simon Fraser University, British Columbia, Canada, 4.0/4.0.**

Senior Supervisor : Dr. Arrvindh Shriraman

My research is directed at facilitating energy efficient computation via specialization. I have adopted a two-pronged approach. First, a top down approach uses program analysis to determine program regions amenable for specialization using LLVM. Second, a bottom up approach evaluated architectural specialization to enable the efficient offload of accelerated program regions.

The former *workload-first* approach, uses program analysis to analyse and reconstruct program regions to aid the design and evaluation of specialized accelerators. An analysis of twenty-nine workloads revealed significant merit in analysis at the path granularity for specialization (IISWC'16). Further analysis of instruction dependency chains in frequently executed paths revealed opportunities for specialized macro-instructions (MICRO'16). An insight into the nature of frequently occurring paths led to the development of a new program abstraction for accelerators (HPCA'17). Robust alias analysis at the path granularity also enabled low overhead memory access interfaces for accelerators (submitted ASPLOS'17). I am also leading an ongoing effort to transparently generate application binaries with specialized regions offloaded to a tightly coupled FPGA substrate.

For the latter *architecture-first* approach, I designed and evaluated a hardware accelerator for software data structures. The access of and compute on data structures is offloaded to an array of processing elements which are tightly coupled to the last level cache (ICS'15). I also evaluated a specialized coherence protocol for fixed function accelerators (ISCA'15) which improves performance and reduces energy consumption by mitigating redundant data movement.

Publications : **HPCA'17, IISWC'16, MICRO'16, ICS'16, ISCA'15, ICS'15**

01/11 – 04/13 **MSc in Computing Science, Simon Fraser University, British Columbia, Canada, 3.8/4.0.**

Senior Supervisor : Dr. Arrvindh Shriraman

Designed and evaluated a variable granularity cache memory hierarchy. The system adaptively varies the cache line size to eliminate data fetches not used by the application. Workloads benefited from increased effective cache space. Overall cache miss rates improved and dynamic energy consumption was reduced. The proposed architecture was modeled using the RUBY memory system simulator and evaluated on twenty-two workloads drawn from popular benchmark suites. A subsequent research work evaluated a variable granularity cache coherence protocol.

Publications : **ISCA'13, MICRO'12**

08/06 – 04/10 **B. Tech in Computer Engineering, Biju Patnaik University of Technology, Orissa, India, 8.3/10.0.**

Supervisor : Dr. Satyananda Champati Rai

Designed and implemented a genetic algorithm to address the problem of channel allocation in cellular networks. The algorithm computes a pseudo optimal borrowing scheme amongst neighbouring cells. The implementation used variable separation to reduce the search space. The approach improved over the state of the art and consistently computed near optimal solutions.

Publications

- 2017 **Needle : Leveraging program analysis to extract accelerators from whole programs**, Snehasish Kumar, Nick Sumner, Vijayalakshmi Srinivasan, Steve Margem, and Arrvindh Shriraman, In *Proceedings of the 23rd ACM International Conference on High Performance Computer Architecture*, HPCA 2017.
- 2016 **ChainSaw : Creating Von-Neumann Accelerators with Fused Instruction Chains**, Amirali Sharifian, Snehasish Kumar, Apala Guha, and Arrvindh Shriraman, In *Proceedings of the 49th Annual IEEE/ACM International Symposium on Microarchitecture*, MICRO 2016.

SPEC-AX : Extracting Accelerator Benchmarks from Microprocessor Benchmarks,
Snehasish Kumar, Nick Sumner, and Arrvindh Shriraman, In *Workload Characterization (IISWC), 2016 IEEE International Symposium on*, IISWC 2016.

Peruse and Profit: Estimating the Accelerability of Loops,
Snehasish Kumar, Vijayalakshmi Srinivasan, Amirali Sharifian, Nick Sumner, and Arrvindh Shriraman, In *Proceedings of the 2016 International Conference on Supercomputing*, ICS '16, pages 21:1–21:13.

2015 **Fusion: Design Tradeoffs in Coherent Cache Hierarchies for Accelerators,**
Snehasish Kumar, Arrvindh Shriraman, and Naveen Vedula, In *Proceedings of the 42Nd Annual International Symposium on Computer Architecture*, ISCA '15, pages 733–745.

DASX: Hardware Accelerator for Software Data Structures,
Snehasish Kumar, Naveen Vedula, Arrvindh Shriraman, and Vijayalakshmi Srinivasan, In *Proceedings of the 29th ACM on International Conference on Supercomputing*, ICS '15, pages 361–372.

2013 **Protozoa : Adaptive Granularity Cache Coherence,**
Hongzhou Zhao, Arrvindh Shriraman, Snehasish Kumar, and Sandhya Dwarkadas, In *Proceedings of the 40th Annual International Symposium on Computer Architecture*, ISCA 2013.

Architectural Support for a Variable Granularity Cache Memory System,
Snehasish Kumar, *MSc Thesis*.

2012 **Amoeba-Cache : Adaptive Blocks for Eliminating Waste in the Memory Hierarchy,**
Snehasish Kumar, Hongzhou Zhao, Arrvindh Shriraman, Eric Matthews, Sandhya Dwarkadas, and Lesley Shannon, In *Proceedings of the 45th Annual IEEE/ACM International Symposium on Microarchitecture*, MICRO 2012.

Workshops, Posters & Presentations

- 01/16 SRC India Design Review, Intel Bangalore – CoolCaches
- 06/15 SFU-ZU workshop on Big Data – Data Structure Accelerators
- 12/13, 08/14 WoNDP'13, PACT'14 – SQRL : Hardware Accelerator for Collecting Software Data Structures

Awards

- 08/16 President's PhD Scholarship, Simon Fraser University
- '16, '14, '12 Graduate Fellowship, Simon Fraser University
- 01/14 Special Graduate Entrance Scholarship, Simon Fraser University

Projects

- 01/15 Networks : Parallel implementation of Kou, Markowsky and Berman (1981) algorithm
- 04/14 Natural Language Processing : Optimizing the Bitpar CKY parser
- 12/11 Computational Geometry : Interactive demo for the Linear Cell Complex (CGAL)
- 04/11 Machine Learning : Non-Negative Matrix Factorisation for large datasets

Professional and Academic Experience

- 06/13 – 12/13 Research Intern : Systems Technology and Architecture
IBM, T.J. Watson Research Centre
- '11 – '16 Research Assistant : SYNAR Group, Simon Fraser University
- '11, '13 Teaching Assistant : CMPT 880, 120, 165, 300

Skills

Languages C++11, C, Python

Frameworks LLVM Compiler Infrastructure, Intel PIN

Simulators Multifacet GEMS (Ruby), MacSim

Leadership

05/12 – 03/13 Councillor, Graduate Student Society, Simon Fraser University

11/08 – 07/10 Microsoft Student Partner

09/09 – 07/10 Treasurer, IEEE Student Chapter

Languages

English Proficient

Hindi Proficient

Bengali Spoken