

List of publications

Snehasish Kumar

November 13, 2016

1. Snehasish Kumar, Nick Sumner, Vijayalakshmi Srinivasan, Steve Margerm, and Arrvindh Shriraman. Needle : Leveraging program analysis to extract accelerators from whole programs. In *23rd ACM International Conference on High Performance Computer Architecture*, HPCA '17, feb 2017. 22%.
2. Amirali Sharifian, Snehasish Kumar, Apala Guha, and Arrvindh Shriraman. ChainSaw : Creating Von-Neumann Accelerators with Fused Instruction Chains. In *49th Annual IEEE/ACM International Symposium on Microarchitecture*, MICRO '16, Washington, DC, USA, oct 2016. IEEE Computer Society. 22%.
3. Snehasish Kumar, Nick Sumner, and Arrvindh Shriraman. SPEC-AX : Extracting Accelerator Benchmarks from Microprocessor Benchmarks. In *2016 IEEE International Symposium on Workload Characterization*, IISWC '16, pages 1–11, September 2016. 30%.
4. Snehasish Kumar, Vijayalakshmi Srinivasan, Amirali Sharifian, Nick Sumner, and Arrvindh Shriraman. Peruse and Profit: Estimating the Accelerability of Loops. In *30th ACM International Conference on Supercomputing*, ICS '16, pages 21:1–21:13, New York, NY, USA, 2016. ACM. 24%.
5. Snehasish Kumar, Arrvindh Shriraman, and Naveen Vedula. Fusion: Design Tradeoffs in Coherent Cache Hierarchies for Accelerators. In *42nd Annual International Symposium on Computer Architecture*, ISCA '15, pages 733–745, New York, NY, USA, 2015. ACM. 19%.
6. Snehasish Kumar, Naveen Vedula, Arrvindh Shriraman, and Vijayalakshmi Srinivasan. DASX: Hardware Accelerator for Software Data Structures. In *29th ACM International Conference on Supercomputing*, ICS '15, pages 361–372, New York, NY, USA, 2015. ACM. 25%.
7. Hongzhou Zhao, Arrvindh Shriraman, Snehasish Kumar, and Sandhya Dwarkadas. Protozoa : Adaptive Granularity Cache Coherence. In *40th Annual International Symposium on Computer Architecture*, ISCA '13, New York, NY, USA, jun 2013. ACM. 19%.
8. Snehasish Kumar, Hongzhou Zhao, Arrvindh Shriraman, Eric Matthews, Sandhya Dwarkadas, and Lesley Shannon. Amoeba-Cache : Adaptive Blocks for Eliminating Waste in the Memory Hierarchy. In *45th Annual IEEE/ACM International Symposium on Microarchitecture*, MICRO '12, Washington, DC, USA, dec 2012. IEEE Computer Society. 18%.